



ML Deployment Workshop Plan: From Training to Model Drift

Duration: 3-4 hours

Audience: Engineers new to ML

Format: Live coding + guided lab (EC2 and Colab friendly)



Core Philosophy

1. **Start Simple, Build Intuition:** Don't overwhelm with tooling. Use the minimal working pipeline and scale concepts later.
2. **Everything Breaks in Prod:** Emphasize drift, monitoring, and data mismatches after deployment.
3. **Local-to-Cloud Bridge:** Most real ML starts locally but lives on the cloud. Teach reproducibility and deployment as a habit.



Workshop Breakdown



Part 1: Setup (15 min)

Objective: Ensure everyone is ready to run notebooks and SSH into EC2

- ☒ Google Colab (or local Python) for training
- ☒ AWS EC2 instance (t2.medium or similar) pre-created (SSH via PEM)
- ☒ Install Python 3.10+, pip install fastapi uvicorn scikit-learn numpy matplotlib
- ☒ GitHub repo with starter code



Part 2: Train a Simple Classifier (30 min)

Objective: Basic ML training loop using scikit-learn or PyTorch

- Dataset: CIFAR-10 or MNIST (download via code)
- Model: Simple CNN or sklearn RandomForestClassifier
- Metrics: Accuracy on train/test split
- Save model using joblib or torch.save



Deliverable:

`saved_models/
model.joblib`

Part 3: Build a FastAPI Inference Server (30 min)

Objective: Serve the model over an API

- Write a FastAPI app:
 - /ping → health check
 - /predict → takes an image (base64 or URL), returns class
- Test locally on Colab (via ngrok or mock request)
- Package as app.py + requirements.txt

 Deliverable:

```
fastapi_app/  
app.py  
requirements.txt  
saved_models/model.joblib
```

Part 4: Deploy to EC2 (30 min)

Objective: Deploy inference server to EC2

- SCP code to EC2
- SSH in, install dependencies
- Run FastAPI with `uvicorn app:app --host 0.0.0.0 --port 8000`
- Open port 8000 in EC2 security group
- Call prediction endpoint from Colab/local





Part 5: Simulate Model Drift (30 min)

Objective: Show how model performance can degrade post-deployment

- Use CIFAR-10 → Classify CIFAR-100 images (wrong domain)
- Log accuracy drops
- Optionally: Add versioning (`model_v1.joblib`, `model_v2.joblib`)
- Teach: Need for retraining pipelines, monitoring

Bonus: Real-World Gotchas (15 min)

Objective: Highlight challenges post-deployment

-  Inference mismatches (different image preproc)
-  Performance drops with noisy inputs
-  Versioning models & rollback
-  Live metrics tracking (demo with a simple CSV or plot)